

Transactional Agent Model for Fault-Tolerant Object Systems

Tomoaki Kaneda, Youhei Tanaka, Tomoya Enokido, and Makoto Takizawa
 Dept. of Computers and Systems Engineering
 Tokyo Denki University, Japan
 {kaneda, youhei, eno, taki}@takilab.k.dendai.ac.jp

ABSTRACT

A transactional agent is a mobile agent which manipulates objects in multiple computers by autonomously finding a way to visit the computers. The transactional agent commits only if its commitment condition like atomicity is satisfied in presence of faults of computers. On leaving a computer, an agent creates a surrogate agent which holds objects manipulated. A surrogate can recreate a new incarnation of the agent if the agent itself is faulty. If a destination computer is faulty, the transactional agent finds another operational computer to visit. After visiting computers, a transactional agent makes a decision on commitment according to its commitment condition. We discuss design and implementation of the transactional agent which is tolerant of computer faults.

Categories and Subject Descriptors

H.2.4 [Systems]: Transaction processing

General Terms

Algorithms, Reliability

Keywords

Mobile agent, Transaction, Fault-Tolerant

1. INTRODUCTION

A transaction manipulates multiple objects distributed in computers through methods. Objects are encapsulations of data and methods for manipulating the data. A transaction is modeled to be a sequence of methods which satisfies the ACID (atomicity, consistency, isolation, and durability) properties [8, 9]. Huge number and various types of peer computers are interconnected in peer-to-peer (P2P) networks [3]. Personal computers easily get faulty not only by crash but also by hackers and intrusions. A mobile agent can autonomously escape from faulty computers by moving

to another operational computer. Mobile agents [5, 19] are programs which move to remote computers and then locally manipulate objects on the computers.

An ACID transaction initiates a subtransaction on each database server, which is realized in mobile agents [16, 9, 13]. In this paper, a *transactional agent* is a mobile agent which autonomously decides in which order the agent visits computers in presence of computer faults, and locally manipulates objects in a current computer with not only atomicity but also other types of commitment conditions like at-least-one condition [6]. After manipulating all or some objects in computers, an agent makes a decision on *commit* or *abort*. For example, an agent atomically commits only if all objects in the computers are successfully manipulated [4]. An agent commits if objects in at least one of the computers are successfully manipulated. In addition, an agent negotiates with another agent which would like to manipulate a same object in a conflicting manner. Through the negotiation, each agent autonomously makes a decision on whether the agent holds or releases the objects [6, 14].

If an agent leaves a computer, objects locked by the agent are automatically released. Hence, once leaving a computer, an agent cannot abort. An agent creates a *surrogate agent* on leaving a computer. A surrogate agent still holds locks on objects in a computer on behalf of the agent after the agent leaves.

A transactional agent autonomously finds another destination computer if a destination computer is faulty. An agent and surrogate are faulty if the current computer is faulty. Some surrogate of the agent which exists on another computer recreates a new incarnation of the agent. Similarly, if a surrogate may be faulty, another surrogate detects the fault and takes a way to recover from the fault. For example, if an agent takes an at least one commitment condition, a fault of the surrogate can be neglected as long as at-least-one surrogate is operational.

In section 2, we present a system model. In section 3, we discuss transactional agents. In section 4, we discuss fault-tolerant mechanism. In sections 5 and 6, we discuss implementation and evaluation of transactional agents.

2. SYSTEM MODEL

A system is composed of *computers* interconnected in reliable networks. Each computer is equipped with a class base (*CB*) where classes are stored and an *object base* (*OB*) which is a collection of persistent objects. A *class* is composed of attributes and methods. An object is an instantiation of a class which is an encapsulation of data and meth-

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

SAC'05 March 13-17, 2005, Santa Fe, New Mexico, USA
 Copyright 2005 ACM 1-58113-964-0/05/0003 ...\$5.00.

ods. If result obtained by performing a pair of methods op_1 and op_2 on an object depends on the computation order, op_1 and op_2 *conflict* with one another. For example, a pair of methods *increment* and *reset* conflict on a *counter* object. On the other hand, *increment* and *decrement* do not conflict, i.e. are *compatible*.

A transaction is modeled to be a sequence of methods, which satisfies the ACID properties [4]. Especially, a transaction can commit only if all the objects are successfully manipulated. If a method op_1 from a transaction T_1 is performed before a method op_2 from another transaction T_2 which conflicts with op_1 , every method op_3 from T_1 has to be performed before every method op_4 from T_2 conflicting with the method op_3 . This is the *serializability* property [2, 4]. Locking protocols [2, 4, 7] are used to realize the serializability of transactions. Here, a transaction locks an object before manipulating the object.

A *mobile agent* is a program which moves around computers and locally manipulates objects in each computer [5, 18, 19]. A mobile agent is composed of classes. A home computer $home(c)$ of a class c is a computer where the class c is stored. For example, each class c is identified by a pair of IP address of a home computer $home(c)$ and a local path to the directory where the class c is stored. A home computer $home(A)$ of a mobile agent A is a home computer of the class of the agent A .

3. TRANSACTIONAL AGENTS

3.1 Model of transactional agent

A *transactional agent* is a mobile agent which satisfies the following properties:

1. autonomously decides on which computer to visit.
2. manipulates objects on multiple computer.
3. commits only if some commitment condition of the agent is satisfied, otherwise aborts.

For simplicity, a term *agent* means a transactional agent in this paper. *Target* objects are objects to be manipulated by an agent. *Target* computers have the target objects. An agent A is composed of *routing* $RC(A)$, *commitment* $CC(A)$, and *manipulation* agents $MC(A, D_1), \dots, MC(A, D_n)$, where D_i stands for a target computer of the agent A . Here, let $Dom(A)$ be a set of *target* computers D_1, \dots, D_n of an agent A . First, an agent A on a current computer has to move to a computer in $Dom(A)$. A computer D_j to which an agent A on D_i moves is a *destination* computer. An agent A has to autonomously make a decision on which computer to visit. In the routing agent $RC(A)$, a destination computer is selected. Then, the agent A moves to the destination computer. Here, an agent first finds a candidate set of possible destination computers. Then, the agent selects one target computer in the candidate computers and moves to the computer.

Secondly, a transactional agent A manipulates objects in a current computer D . The agent A initiates a manipulation agent $MC(A, D)$ for manipulating objects in the current computer D from the home computer. If an object base is realized in a relational database system [11], objects are manipulated by issuing SQL commands in $MC(A, D)$.

Lastly, a transactional agent makes a decision on whether the agent can commit or abort after visiting target computers. A traditional transaction [2] *atomically* commits only if

objects in all the target computers are successfully manipulated. In this paper, we consider other types of commitment conditions [6]. For example, in the at-least-one commitment, a transaction can commit only if objects in at least one target computer are successfully manipulated.

3.2 Routing agent

A transactional agent A locally manipulates objects in a computer D_i through the manipulation agent $MC(A, D_i)$ and then outputs intermediate objects $OUT(A, D_i)$. In the meanwhile, the agent A visits another computer D_j . Here, objects in D_j are manipulated through the manipulation agent $MC(A, D_j)$ by using the intermediate objects $In(A, D_j)$ ($=OUT(A, D_i)$). Thus, the manipulation classes are related with input-output relation. Here, $D_i \xrightarrow{x} D_j$ shows that the manipulation agent $MC(A, D_i)$ outputs an intermediate object x which is used by $MC(A, D_j)$. If $D_i \xrightarrow{x} D_j$, the agent A has to visit D_i before D_j and the intermediate object x has to be delivered to D_j . The input-output relation is shown in an input-output graph as shown in Figure 1.




Figure 1: Input-output graph

There are *computer* and *object* nodes. Directed edges $D_i \rightarrow x$ and $x \rightarrow D_i$ show that the manipulation agent $MC(A, D_i)$ outputs and inputs an object x , respectively. In Figure 1, the agent A outputs an intermediate object w in D_1 . The agent A uses x in D_3 , D_4 , and D_5 . This means the agent A is required to visit D_3 , D_4 , and D_5 after D_1 .

From the input-output graph, a transactional agent A decides in which order the agent visits. A directed acyclic graph (DAG) $Map(A)$ named a *map* is created from the input-output graph [Figure 2]. Here, a node D shows a computer D with a manipulation agent $MC(A, D)$. A directed edge $D_1 \rightarrow D_2$ a computer D_2 is required to be manipulated after D_1 . $D_1 \rightarrow^* D_2$ if and only if (iff) $D_1 \rightarrow D_2$ or $D_1 \rightarrow D_3 \rightarrow^* D_2$ for some computer D_3 . D_1 and D_2 are independent ($D_1 \parallel D_2$) if neither $D_1 \rightarrow^* D_2$ nor $D_2 \rightarrow^* D_1$. Here, a transactional agent A can visit the computers D_1 and D_2 in any order and can in parallel visit the computers D_1 and D_2 . Figure 2 shows an example of a map $Map(A)$ obtained from the input-output graph of Figure 1. Here, an agent A is required to visit a computer D_3 after D_1 , D_4 after D_2 and D_3 , and D_5 after D_4 . On the other hand, an agent A can visit D_1 and D_2 in any order, even in parallel.

In Figure 1, the intermediate object w has to be delivered to D_3 , D_4 , and D_5 . There are following ways to bring an intermediate object x obtained in D_i to D_j :

1. A transactional agent A carries the intermediate object x to D_j .
2. x is transferred from D_i to D_j before A arrives at D_j .
3. x is transferred from D_i to D_j after A arrives at D_j .

A routing agent $RC(A)$ of a transactional agent A with a map $Map(A)$ is moving around computers [Figure 3]. First,




Figure 2: Map.

a collection I of computers which do not have any in-coming edge are found in $Map(A)$. For example, $I = \{D_1, D_2\}$ in Figure 2. One computer D_i is selected in I so as to satisfy some condition, e.g. D_i nearest to the current computer is selected. For example, an agent takes a computer D_1 in Figure 2. The agent A moves to D_i . Here, a manipulation agent $MC(A, D_i)$ is loaded to D_i from the home computer. After manipulating objects in D_i , D_i is removed from $Map(A)$. Another destination D_j is selected and A moves to D_j .

Initially, a routing agent $RC(A)$ of the agent A is loaded and started on a computer. The computer is a *base* computer $base(A)$ of the agent A . An agent A leaves the base computer for a computer D_i . Here, D_i is a *current* computer $current(A)$ of A . If the agent A invokes a method t of a class c on D_i , the class c is searched:

1. The cache of the current computer D_i is first searched for the class c . If c is found in the cache, the method t in the cache is invoked.
2. If not, the class base (CB_i) of D_i is locally searched. If found, the class c in CB_i is taken to invoke t .
3. Otherwise, the class c is transferred from the home computer $home(c)$ into D_i .

A history $H(A)$ shows a sequence of computers which an agent A has visited.




Figure 3: Mobile agent.

3.3 Manipulation agent

A manipulation agent is composed of not only application-specific classes but also library classes like JDBC [17] and JAVA classes [18]. Each computer is assumed to support a platform to perform a mobile agent on an object base (OB). A platform includes *cache* and *class base* (CB). The routing, manipulation, and commitment agents of a transactional agent A are stored in the class base (CB) of the home computer $home(A)$. If an agent A invokes a method t of a class c in a computer D_i , the class c is loaded from the home computer $home(c)$ to the cache in D_i . Then, the method t of the class c is performed in D_i . If a method u of another class d is invoked in the method t , the class d is loaded from the home computer $home(d)$ as well as the class c . Meanwhile, if another agent B invokes a method t of the class c in D_i , the class c in the cache is used to invoke the method t without loading the class c . Thus, if classes are cached in a computer D_i , methods in the classes are locally invoked in D_i without any communication. Otherwise, it takes a longer time to invoke methods since classes with the

methods are transferred from the home computers in networks. Here, the class c is loaded i.e. *cached* to D_i . The method t of the class c is performed on D_i . If another agent B comes to D_i after A has left D_i , B can take usage of the class c in the cache.

3.4 Commitment agent

If a transactional agent A finishes manipulating objects in each computer, the following *commitment* condition is checked by the commitment agent $CC(A)$:

1. *Atomic commitment*: an agent is successfully performed on all the computers in the domain $Dom(A)$, i.e. all-or-nothing principle used in the traditional two-phase commitment protocol [4, 15].
2. *Majority commitment*: an agent is successfully performed on more than half of the computers in $Dom(A)$.
3. *At-least-one commitment*: an agent is successfully performed on at least one computer in $Dom(A)$.
4. $\binom{n}{r}$ *commitment*: an agent is successfully performed on more than r out of n computers ($r \leq n$) in $Dom(A)$.
5. *Application specific commitment*: condition specified by application is satisfied.

3.5 Resolution of confliction

Suppose an agent A moves to a computer D_j from another computer D_i . The agent A cannot be performed on D_j if there is an agent or surrogate B conflicting with A . Here, the agent A can take one of the following ways:

1. *Wait*: The agent A in the computer D_i *waits* until the agent A can land at a computer D_j .
2. *Escape*: The agent A *finds* another computer D_k which has objects to be manipulated before D_j .
3. *Negotiate*: The agent A *negotiates* with the agent B in D_j . After the negotiation, B releases the objects or aborts.
4. *Abort*: The agent A *aborts*.

Deadlock among agents may occur. If the timer expires, the agent A takes a following way:

1. The agent A retreats to a computer D_j in the history $H(A)$. All surrogates preceding D_j are aborted.
2. Then, the surrogate agent A_j on D_j recreates a new incarnation of the agent A . The agent A finds another destination computer D_h .

The surrogate A_j to which the agent A retreats plays a role of checkpoint [12].

Suppose a surrogate agent B holds an object in a computer D_j . An agent A would like to manipulate the object but conflicts with B in D_j . The surrogate B makes a following decision:

1. *Atomic commitment*: The agent A waits until the surrogate B finishes.
2. *At-least-one commitment*: If the surrogate B knows at least one sibling surrogate of B is committable, B releases the object and aborts after informing the other sibling surrogates of this abort.
3. *Majority commitment*: If the surrogate B knows more than half of the sibling surrogates are committable, B releases the object and aborts after informing the other surrogates.
4. $\binom{n}{r}$ *commitment*: If the surrogate B knows more than

or equal to r sibling surrogate agents are committable, the surrogate B releases the object and aborts.

4. FAULT-TOLERANT AGENT

We assume computers may stop by fault and networks are reliable. A transactional agent is faulty only if a current computer of the agent is faulty. Suppose an agent A finishes manipulating objects on a computer D_i . The agent A selects one computer D_j from the map $Map(A, D_i)$. The agent A detects by timeout mechanism that D_j is faulty. The agent A tries to find another destination computer D_k [Figure 4]. If found, A moves to D_k as presented here. If A cannot find another destination computer in $Map(A, D_i)$, the agent A backs to the preceding computer D_k [Figure 5]. D_i is removed from $Map(A, D_k)$. Then, the agent in D_k tries to find another destination computer in $Map(A, D_k)$.




Figure 4: Forwarding recovery.




Figure 5: Backwarding recovery.

An agent A leaves its surrogate agent A_i on a computer D_i . The surrogate A_i holds objects even after the agent A leaves D_i . An agent A and surrogate agent A_i stop if the current computers are faulty. First, suppose an agent A stops on the current computer D_j . Suppose that the agent A comes from D_i to D_j . The surrogate A_i on D_i detects that the agent A stops on D_j . Here, A_i takes one of the following actions:

1. Find a succeeding surrogate A_k of A_i and skips A_j .
2. Recreate a new incarnation of the agent A .

If the commitment condition is not atomic, the surrogate A_j takes the first one, i.e. skips the fault of A_j . For the atomic condition, A_i recreates a new incarnation of the agent A . The agent A takes another destination computer D_k in $Map(A, D_i)$. If found, the agent A moves to D_k . Otherwise, A waits until the computer D_i is recovered or backs to the precedent computer from D_j .

A surrogate A_i on a computer D_i may be faulty as well. A preceding surrogate A_j on D_j detects the fault of A_i . Suppose a surrogate agent A_i of A exists on D_i . A_{i+1} and A_{i-1} show the succeeding and preceding surrogate agents of A_i , respectively [Figure 6]. A_i periodically sends an enquiry message *AYL* (are you alive) to A_{i+1} and A_{i-1} to check if A_{i+1} and A_i are alive. On receipt of the *AYL* message, a surrogate sends back a response message *IAL* (I am alive). Thus, a faulty surrogate is detected by the succeeding and preceding a surrogate with timeout mechanism.

If A_i detects the stop of A_{i+1} , A_i does the followings:

1. A new incarnation of the agent A is recreated on D_i .
2. From the map $Map(A, D_i)$, a new destination D different from D_{i-1} is detected.




Figure 6: Fault detection.

3. If detected, the agent A moves to D . Otherwise, A_i informs A_{i-1} of *abort* and then aborts. A_{i-1} does the procedure from step 1.

If the surrogate A_i detects the stop of the preceding surrogate A_{i-1} or receives an *abort* message for A_{i-1} , A_i informs the succeeding surrogate A_{i+1} of *abort*. On receipt of the *abort* message from A_i , A_{i+1} forwards the *abort* message to A_{i+2} and then aborts. Thus, *abort* messages are eventually forwarded up to the agent A . In Figure 7, suppose A_2 stops. A pair of surrogates A_1 and A_3 detect the stop of A_2 . A_1 creates a new incarnation A' of the agent A . The obsolete incarnation A still is moving to D_6 . The succeeding surrogate A_3 of A_2 sends an *abort* message to A_4 . If the *abort* message catches up the agent A , A can be aborted. Otherwise, the obsolete incarnation A cannot stop. Thus, there might exist multiple incarnations of an agent.




Figure 7: Incarnations of an agent.

On receipt of an *AYL* message from the preceding surrogate A_{i-1} , A_i sends an *IAL* message with the address information which A_i knows of surrogates are backwarding to preceding surrogates. If the surrogate A_i finds A_{i-1} to be faulty, A_i sends an *abort* message to not only A_{i+1} but also a surrogate whose address A_i knows and which is nearest to the current computer of A . By this method, an *abort* message can more easily catch up with the agent mapped the agent can be aborted.

5. IMPLEMENTATION

We discuss how to implement transactional agents in Aglets. A transactional agent A is composed of *routing*, *manipulation*, and *commitment* subagents. An routing agent $RC(A)$ with a map $Map(A)$ is transferred from one computer to another. When an agent A , i.e. routing agent $RC(A)$ arrives at a computer D_i , a *manipulation agent* $MC(A, D_i)$ is created by loading the manipulation class.

An object base (*OB*) is realized in a relational database system, Oracle [11]. A transactional agent manipulates *table* objects by issuing SQL commands, i.e. **select**, **insert**, **delete**, and **update** in a current computer D_i . The computation of each agent A on a computer D_i is realized as a local *transaction* on a database system. If the agent A leaves D_i , the transaction for A commits or aborts. That is, objects manipulated by A are released. Even if the agent A leaves D_i , the objects manipulated by A are required to be still held because A may abort after leaving D_i . If the

objects are released, the agent is *unrecoverable*. Therefore, a *surrogate agent* is created on D_i . The surrogate agent is composed of a manipulation agent $MC(A, D_i)$ and an object agent OBA_i . OBA_i behaves as follows:

1. On arrival at a computer D_i , the routing agent $RC(A)$ of an agent A initiates a manipulation agent $MC(A, D_i)$ and an object agent OBA_i on D_i , i.e. $MC(A, D_i)$ and OBA classes are loaded. OBA_i initiates a transaction on an object base OB_i .
2. If $MC(A, D_i)$ issues a method for manipulating objects, OBA_i issues SQL commands to the database system in D_i .
3. If the agent A finishes, A leaves D_i . However, OBA_i is still operational and holding the objects in D_i .
4. OBA_i commits and aborts if the agent A sends *commit* and *abort* requests to the surrogate A_i , respectively.

An object agent OBA_i stays on a computer D_i while holding objects even if the agent A leaves D_i . OBA_i is a local transaction on an object base OB_i . On completion of the agent A , OBA_i and $MC(A, D_i)$ are terminated.




Figure 8: Object agent (OBA).

An OBA class can be loaded to a computer with any type of database system. If a transactional agent comes to D_i from another home computer, an OBA class is loaded to D_i from the home computer. Thus, OBA instances are accumulated in the cache. In order to resolve this problem, an OBA class is loaded as follows:

1. If the OBA class is not cached in the current computer, the OBA class is loaded from $home(OBA)$.
2. If the OBA class could not be loaded from $home(OBA)$, an OBA class in the home computer of the agent is loaded to a computer.

The routing agent $RC(A)$ leaves a computer D_i if the manipulation agent $MC(A, D_i)$ finishes manipulating objects. $MC(A, D_i)$ recreates a new incarnation of a routing agent $RC(A)$ if the agent A stops due to the computer fault.

A transactional agent A can commit if all or some of the surrogates commit depending on the commitment condition. For example, a transactional agent commits if all the surrogate agents successfully exist. Communication among an agent and its surrogate agents is realized by using the XA interface [20] which supports the two-phase commitment protocol [15] [Figure 8]. Each surrogate agent issues a *prepare* request to a computer on receipt of a *prepare* message from A . If *prepare* is successfully performed, the surrogate agent sends a *prepared* message to A . Here, the surrogate agent is *committable*. Otherwise, the surrogate agent aborts after sending *aborted* to A . The agent A receives responses from the surrogate agents after sending *prepare* to the surrogates.

On receipt of the responses from surrogate agents, the agent A makes a decision on *commit* or *abort* based on the commitment condition. For example, if the atomic condition holds, A sends *commit* only if *prepared* is received from every surrogate. The agent A sends *abort* to all committable agents if an *aborted* message is received from at least one surrogate. On receipt of *abort*, a committable surrogate aborts. In the at-least-one commitment condition, A sends *commit* to all committable surrogates only if *prepared* is received from at least one surrogate.

Next, we discuss how to support robustness against faults of computers. Suppose a surrogate agent A_i of a transactional agent A stops after sending *prepared*. Here, A_i is committable. On recovery of the committable surrogate A_i , A_i unilaterally commits if the surrogate agent is committable in the at-least-one commitment condition. In the atomic condition, A_i asks the other surrogates if they had committed. Suppose A_i is abortable, i.e. faulty before receiving *prepared*. On recovery, A_i unilaterally aborts.

6. EVALUATION

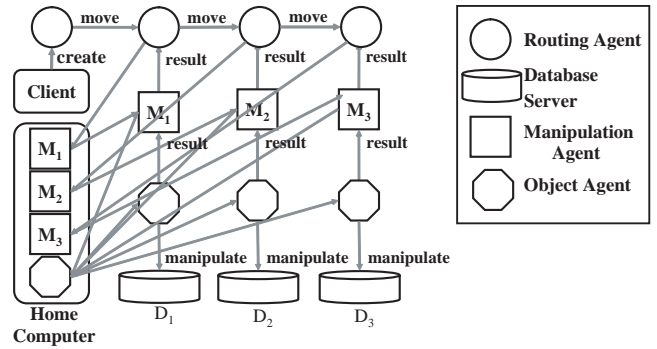


Figure 9: Evaluation model

We evaluate the transactional agent which is implemented in Agetls. In the evaluation, There are three server computers D_1 , D_2 , and D_3 . A transactional agent is created in a computer C by loading classes from the home computer h . The servers D_1 , D_2 , and D_3 are realized in personal computers (Pentium 3) with Oracle database systems, which are interconnected in the 1Gbps Ethernet.

First, a transactional agent A is initiated in a base computer C . The agent A finds in which order D_1 , D_2 , and D_3 to be visited. Here, the agent A visits D_1 , D_2 , and D_3 in this order as shown in Figure 9. On arrival of the agent A on D_i , the manipulation agent $MC(A, D_i)$ and object agent OBA_i are loaded to D_i [Figure 9].

We consider that following types of transactional agents:

- A. The manipulation agents $MC(A, D_1)$ derives intermediate object I from the object base. The object bases in D_2 and D_3 are updated by using the object I , i.e. objects in I are added to the object base.
- B. $MC(A, D_1)$ and $MC(A, D_2)$ derive objects to intermediate objects I_1 and I_2 , respectively. Then, the object base in D_3 is manipulated by using I_1 and I_2 .

There are three ways to deliver intermediate objects derived to another computer:

1. The transactional agent A carries intermediate objects to a destination computer D_j from D_i .

2. After the agent A arrives at a computer D_j , the agent A requests D_i to send the intermediate objects.
3. The agent A transfers the intermediate object I to a computer D_j before leaving D_i .

The total response time of a transactional agent is measured for number of intermediate objects, i.e. number of tuples derived in computeres. Figures 10 and 11 show the response time for the types of transactional agents A and B, respectively. The second and third ways to deliver intermediate objects to destination computers imply shorter response time than the first way.




Figure 10: Response A




Figure 11: Response B

7. CONCLUDING REMARKS

The authors discussed a transactional agent model to manipulate objects in multiple computers with types of commitment constraints in presence of computer faults. A transactional agent autonomously finds a destination computer, moves to a computer, and then locally manipulates objects. We discussed how to implement transactional agents in Aglets and Oracle. We evaluated the transactional agent in terms of response time.

8. REFERENCES

- [1] American National Standards Institute. *The Database Language SQL*, 1986.
- [2] P. A. Bernstein, V. Hadzilacos, and N. Goodman. *Concurrency Control and Recovery in Database Systems*. Addison-Wesley, 1987.
- [3] L. Gong. *JXTA: A Network Programming Environment*, pages 88–95. IEEE Internet Computing,, 2001.
- [4] J. Gray and A. Reuter. *Transaction Processing : Concepts and Techniques*. Morgan Kaufmann, 1993.
- [5] IBM Corporation. *Aglets Software Development Kit Home*. <http://www.trl.ibm.com/aglets/>.
- [6] T. Komiya, T. Enokido, and M. Takizawa. Mobile agent model for transaction processing on distributed objects. *Information Sciences*, 154:23–38, 2003.
- [7] F. H. Korth. Locking primitives in a database system. *Journal of ACM*, 30(1):55–79, 1989.
- [8] N. A. Lynch, M. Merritt, A. F. W. Weihl, and R. R. Yager. *Atomic Transactions*. Morgan Kaufmann, 1994.
- [9] K. Nagi. *Transactional Agents : Towards a Robust Multi-Agent System*. Springer-Verlag, 2001.
- [10] A. Omicini, F. Zambonelli, M. Klusch, and R. Tolksdorf. *Coordination of Internet Agents*. Springer-Verlag, 2001.
- [11] Oracle Corporation. *Oracle8i Concepts Vol. 1 Release 8.1.5*, 1999.
- [12] R. S. Pamula and P. K. Srimani. Checkpointing strategies for database systems. *Proc. of the 15th Annual Conf. on Computer Science, IEEE Computer Society*, pages 88–97, 1987.
- [13] S. Pleisch. *State of the Art of Mobile Agent Computing - Security, Fault Tolerance, and Transaction Support*. IBM Corporation, 1999.
- [14] M. Shiraishi, T. Enokido, and M. Takizawa. Fault-tolerant mobile agent in distributed objects systems. *Proc. of the 9th IEEE International Workshop on Future Trends of Distributed Computing Systems (FTDCS 2003)*, pages 145–151, 2003.
- [15] D. Skeen. Nonblocking commitment protocols. *Proc. of ACM SIGMOD*, pages 133–147, 1982.
- [16] A. D. Stefano, L. L. Bello, and C. Santoro. A distributed heterogeneous database system based on mobile agents. *Proc. of the 7th Workshop on Enabling Technologies (WETICE'98)*, IEEE Computer Society, pages 223–229, 1998.
- [17] Sun Microsystems Inc. *JDBC Data Access API*. <http://java.sun.com/products/jdbc/>.
- [18] Sun Microsystems Inc. *The Source for Java (TM) Technology*. <http://java.sun.com/>.
- [19] J. E. White. *Telescript Technology : The Foundation for the Electronic Marketplace*. General Magic Inc., 1994.
- [20] X/Open Company Ltd. *X/Open CAE Specification Distributed Transaction Processing: The XA Specification*, 1991.